


Presentation Strategy & Script: VendorLens

To win over both the **CIO** (who cares about architecture, security, cost, and platform potential) and the **non-technical stakeholders** (who care about usability, speed, and business value), you want to avoid a dry, flat walkthrough.

Instead, frame the presentation as a **dynamic story** centered around a real problem they understand, utilizing the 3-minute processing time to explain the underlying technology.

1. The Presentation Structure & Pacing



2. Step-by-Step Script

Step 1: The Hook (2 Minutes)

- **The Goal:** Establish immediate relevance.
- **What to say:**

"In December 2025, our committee spent weeks manually reviewing complex proposals for the Enterprise LMS RFP. We read hundreds of pages, held long scorecard alignment meetings, and

manually compiled risk reviews.

What if we could automate that first pass? What if we could extract terms, flag policy violations, score them against our actual rubrics, and write a recommendation memo in under three minutes?

That is VendorLens. It's not a hypothetical tool—it is built for the exact procurement we ran."

Step 2: Start the Live Run (30 Seconds)

- **The Goal:** Initiate the process so the audience can watch it compile.
- **Action:**
 1. Open the frontend landing page.
 2. Drag and drop the 3 vendor proposals (Blackboard, Canvas, Brightspace) or select the LMS bundle.
 3. Click **Analyze**.
 4. *Crucial:* Use the live processing mode (?processing). The screen will display the **AgentProgressBar** and the scrolling **ThinkingLog**.
- **What to say:**

"Let's kick off a live analysis run. I'm uploading three simulated vendor responses. I'll hit Analyze now.

While the system processes these documents in the background over the next few minutes, let me explain exactly what is happening under the hood."

Step 3: Explaining the Tech (3 Minutes - While the Pipeline Runs)

Use this time to explain the complex AI architecture using three concrete analogies:

A. Multi-Agent System: "The Specialized Team"

- **The Tech:** LangGraph orchestrating specialized nodes (Extract → Risk → Score → Memo → Negotiation).
- **The Analogy:**

"Look at the progress bar. Right now, it's running a team of specialized AI agents. Instead of asking a single generalist to read the document and write a report in one go, we hired a specialized committee:

 - **Agent 1 (The Researcher):** Extracts key facts from the text.
 - **Agent 2 (The Compliance Officer):** Looks at the facts and compares them to our legal policies.
 - **Agent 3 (The Accountant):** Grades the vendor on a strict 0–10 scale.
 - **Agent 4 (The Writer):** Drafts a beautiful recommendation memo.
 - **Agent 5 (The Negotiator):** Prepares a custom playbook for our sourcing team."

B. RAG (Retrieval-Augmented Generation): "The Open-Book Exam"

- **The Tech:** LlamaIndex + local ChromaDB vector store.
- **The Analogy:**

*"Our compliance agent is running right now. How does it know what is in the document? Standard AI models try to answer questions from memory, which leads to guessing. VendorLens does not do this. It works like an **open-book exam**. Before the AI answers a question, it searches our local document database for the exact page and paragraph, reads it, and bases its answer solely on that text."*

C. Prompt Caching: "The Opened Textbook"

- **The Tech:** Anthropic prompt caching with cache pre-warming.
- **The Analogy:**

"Usually, every time you ask an AI a question, it has to read the entire textbook from page one, which is slow and expensive. With **Prompt Caching**, we open the textbook to the rules section once and keep it open. Every vendor we analyze reads from the same open book, reducing our API costs by 90% and speeding up the analysis by 85%."

Step 4: The Reveal & Deep Dive (3 Minutes)

- **The Goal:** Walk through the completed results.
- **Action:** The pipeline completes, confetti bursts, and the results cards load.
- **What to say:**

"And we are done. The pipeline has successfully analyzed all three vendors."

Notice how the cards sorted themselves automatically: Canvas scored the highest at 8.2, and Blackboard is flagged with a high liability risk."

- **Action:** Hover over the red risk flags to show the tooltips, expand the contract details, and scroll down to the markdown memo. Point out the **Negotiation Playbook** showing the custom BATNA and concession strategies for each vendor.
- **Action:** Click the download button to show how easily the final memo converts to a local file.

3. The CIO Value Pitch (2 Minutes)

Address the CIO's three main concerns directly:

1. Security & Privacy:

- *Message:* "Our vendor proposals never train public models. The vector database (ChromaDB) runs entirely locally in our environment. Analysis jobs live in-memory with a strict 10-minute eviction policy (TTL) so that raw documents are never stored long-term."

2. Platform Flexibility (Swappable Context Bundles):

- *Message:* "VendorLens is a platform, not a single-purpose tool. All department-specific logic—legal policies, evaluation criteria, and rubrics—is decoupled from the code. If we want to run a cybersecurity RFP next month, we swap three text files in the `context_bundles/` folder. The agents, graph, and UI remain completely unchanged."

3. Auditability (LangSmith):

- *Message:* "Every single decision, document retrieval chunk, and LLM call is fully traceable via LangSmith. Sourcing teams can audit exactly why a vendor was given a specific score or why a risk flag was raised."

4. Q&A Cheat Sheet

Question	Answer
"Could we use this for other departments?"	"Yes. We just add a new folder to <code>data/context_bundles/</code> containing their policy and rubric. No code changes needed."
"What does it cost to run?"	"Because we use Claude Haiku for mechanical tasks (scoring, negotiation) and implemented Prompt Caching , a complete run analyzing three major vendor proposals costs less than \$0.05 in API tokens."

"What if the vendor files are PDFs?"	"The ingestion pipeline converts and chunks both .txt and .pdf files automatically."
"Who maintains the database?"	"The vector database (ChromaDB) runs locally in the Docker container, requiring zero database administrator overhead. Completed job records are persisted to a lightweight Postgres instance."